

# Data Wrangling Strategies with R

Clay Ford

Fall 2021

# Agenda

Cover the following topics with examples and exercises:

- ▶ Understanding lists
- ▶ Working with dates
- ▶ Working with character data
- ▶ Binding data frames
- ▶ Merging or Joining data frames
- ▶ Reshaping data frames

## Data structures in R: vector

Vector: 1D object of same data type (eg, all numeric, or all character). Like a column of data in a spreadsheet

```
x <- c(1, 4, 7, 11)  
x
```

```
## [1]  1  4  7 11
```

## Data structures in R: data frame

data frame: 2D object consisting of vectors of the same length but of potentially different data types

```
x <- c(1, 3, 5)
g <- c("M", "M", "F")
df <- data.frame(x, g)
df
```

```
##    x g
## 1 1 M
## 2 3 M
## 3 5 F
```

Notice x is numeric while g is character.

## Data structures in R: lists

list: the most general data structure. It can contain vectors, data.frames, and other lists. When data wrangling, we sometimes need to work with lists at an intermediate step.

```
lst <- list(x = 5, df = data.frame(x, g))  
lst
```

```
## $x  
## [1] 5  
##  
## $df  
##   x g  
## 1 1 M  
## 2 3 M  
## 3 5 F
```

## Applying functions to elements in a list

The `lapply` function allows us to *apply* functions to elements of a list. Example: Find the mean of each list element

```
(lst <- list(y1 = 1:4, y2 = 6:10))
```

```
## $y1  
## [1] 1 2 3 4  
##  
## $y2  
## [1] 6 7 8 9 10
```

```
lapply(lst, mean)
```

```
## $y1  
## [1] 2.5  
##  
## $y2  
## [1] 8
```

## Dates and Date-times in R

- ▶ When a date (eg, April 5, 1982) is formatted as a “date” in R, it becomes the number of days after (or before) January 1, 1970.
- ▶ When a date-time (eg, April 5, 1982 1:15 PM) is formatted as a “date-time”, it becomes the number of seconds after (or before) January 1, 1970.
- ▶ This simplifies the calculation of time spans.
- ▶ The `lubridate` package helps us parse dates and date-times as well as perform calculations and conversions of such values.

## Parsing dates

- ▶ lubridate provides a series of functions for parsing dates that are a permutation of the letters “m”, “d” and “y” to represent the ordering of month, day and year.
- ▶ lubridate provides functions for every permutation of “m”, “d”, “y”.

```
library(lubridate)
d <- "April 5, 1982"
d <- mdy(d)
d
```

```
## [1] "1982-04-05"
```



## Printed values vs. stored values

- ▶ A date that is parsed with `lubridate` will print to the console and appear in data frames as if it's character data.
- ▶ Use `as.numeric` to see the stored value.

```
d
```

```
## [1] "1982-04-05"
```

```
as.numeric(d) # days since 1/1/1970
```

```
## [1] 4477
```

## Parsing date-times

To parse date-times, append either `_h`, `_hm`, or `_hms` to the “mdy” function.

```
d <- "April 5, 1982 1:15 PM"  
d <- mdy_hm(d)  
d
```

```
## [1] "1982-04-05 13:15:00 UTC"
```

```
as.numeric(d) # seconds since 1/1/1970
```

```
## [1] 386860500
```

## Parsing times

lubridate also allows us to parse hours, minutes and seconds using `hm`, `ms` and `hms`

```
t <- c("1:23", "2:34")  
hm(t)
```

```
## [1] "1H 23M 0S" "2H 34M 0S"
```

```
ms(t)
```

```
## [1] "1M 23S" "2M 34S"
```

The output is nicely formatted, but these are stored as seconds

## Extracting date components

lubridate provides functions such as month, day, wday, yday, etc. to extract date components

```
d <- mdy("5/5/01")  
month(d, label = TRUE) # month
```

```
## [1] May
```

```
## 12 Levels: Jan < Feb < Mar < Apr < May < Jun < Jul < Aug
```

```
wday(d, label = TRUE) # week day
```

```
## [1] Sat
```

```
## Levels: Sun < Mon < Tue < Wed < Thu < Fri < Sat
```

# Character strings

Our data often include character strings such as names, locations, descriptions, categories or unwanted “junk”. Examples of manipulating character strings include...

- ▶ Converting text to UPPERCASE or lowercase
- ▶ Extract parts of a string (Eg, extract “23” from “<b>23</b>”)
- ▶ Padding strings with zeroes, so 9, 10, 11 become 009, 010, 011
- ▶ Identify patterns of text for purpose of extracting, replacing or subsetting data

We will use the `stringr` package to work with character strings.

## Character data in R

- ▶ Character data have quotes when printed to console
- ▶ But data with quotes does not mean it's character!
- ▶ use `is.character()` to find out.
- ▶ Character data need to be surrounded with quotes (either single or double) when used in R code

```
(x <- c("a","b","c","12"))
```

```
## [1] "a"  "b"  "c"  "12"
```

```
is.character(x)
```

```
## [1] TRUE
```

## Character versus factor

- ▶ Sometimes data that appear to be character are actually stored as a *factor*
- ▶ factors are character data that are stored as integers but have character labels
- ▶ factors are good for using character data in statistical modeling (eg, ANOVA, regression, etc)
- ▶ If your character data is stored as a factor, R automatically handles conversion to dummy matrices necessary for statistical modeling routines
- ▶ factors do not have quotes when printed to console

## Factor data in R

```
(y <- factor(c("a","b","c","c")))
```

```
## [1] a b c c  
## Levels: a b c
```

```
is.character(y)
```

```
## [1] FALSE
```



## When to convert factors to character

- ▶ If you plan to clean or manipulate character data, make sure it's character, not factor.
- ▶ Change factor to character with `as.character` function

```
(y <- factor(c("a","b","c","c")))
```

```
## [1] a b c c  
## Levels: a b c
```

```
(y <- as.character(y))
```

```
## [1] "a" "b" "c" "c"
```

## Convert case of string

- ▶ `str_to_upper`, `str_to_lower`, and `str_to_title` do what you expect

```
library(stringr)  
str_to_upper("day one")
```

```
## [1] "DAY ONE"
```

```
str_to_lower("DAY ONE")
```

```
## [1] "day one"
```

```
str_to_title("day one")
```

```
## [1] "Day One"
```

## find-and-remove within strings

- ▶ `str_remove` finds first occurrence of specified pattern and removes it

```
# find first - and replace with nothing  
str_remove("434-555-1212", "-")
```

```
## [1] "434555-1212"
```

- ▶ `str_remove_all` finds all occurrences of specified pattern and removes it

```
# find all - and replace with nothing  
str_remove_all("434-555-1212", "-")
```

```
## [1] "4345551212"
```

## find-and-replace within strings

- ▶ `str_replace` finds first occurrence of specified pattern and replaces with specified text

```
str_replace("434-555-1212", "-", ".")
```

```
## [1] "434.555-1212"
```

- ▶ `str_replace_all` finds all occurrences of specified pattern and replaces with specified text

```
str_replace_all("434-555-1212", "-", ".")
```

```
## [1] "434.555.1212"
```

## Pad a string with characters

- ▶ `str_pad` will pad a string with characters. This is useful for zip codes or ID numbers.
- ▶ Specify the width, the side of the padding, and what to pad with.

```
(zips <- c(22904, 06443, 01331))
```

```
## [1] 22904 6443 1331
```

```
str_pad(zips, width = 5, side = "left", pad = "0")
```

```
## [1] "22904" "06443" "01331"
```

# Regular Expressions

- ▶ Regular expressions are a language for describing text patterns (eg: email addresses, social security numbers, html tags)
- ▶ A regular expression is usually formed with some combination of *literal characters*, *character classes* and *modifiers*
  - ▶ literal character example: `state` (looking for "state")
  - ▶ character class example: `[0-9]` (any number 0 - 9)
  - ▶ modifier example: `+` (1 or more of whatever it follows)
- ▶ Regular expression example: `state[0-9]+` finds patterns such as `state1`, `state12`, `state99` but not `state`
- ▶ We will cover just the basics today as they work in R

# Character classes

- ▶ `[0-9]`, `[a-z]`, `[A-Z]`
- ▶ Define your own: `[0-3a-g]`, `[AEIOUaeiou]`
- ▶ Predefined character classes
  - ▶ `[:alpha:]` all letters
  - ▶ `[:digit:]` numbers 0 - 9
  - ▶ `[:alnum:]` Alphanumeric characters (alpha and digit)
  - ▶ `[:blank:]` Blank characters: space and tab
  - ▶ `[:lower:]` lowercase letters
  - ▶ `[:upper:]` UPPERCASE letters
  - ▶ `[:punct:]` Punctuation characters
  - ▶ `[:print:]` Printable characters: `[:alnum:]`, `[:punct:]` and space
  - ▶ `[:space:]` Space characters: tab, newline, vertical tab, form feed, carriage return, space

# Modifiers

- ▶ `^` start of string; or negation inside character class
- ▶ `$` end of string
- ▶ `.` any character except new line
- ▶ `*` 0 or more
- ▶ `+` 1 or more
- ▶ `?` 0 or 1
- ▶ `|` or (alternative patterns)
- ▶ `{}` quantifier brackets: exactly `{n}`; at least `{n,}`; between `{n,m}`
- ▶ `()` group patterns together
- ▶ `\` escape character (needs to be escaped itself in R! `\\`)
- ▶ `[]` character class brackets

Note: precede these with a double backslash if you want to treat them as literal characters.



## Basic regular expression examples

- ▶ Remove one or more letters followed by . and space at beginning of string

```
names <- c("Dr. Claibourn", "Mr. Ford", "Ms. Draber")  
str_remove(names, pattern = "^[:alpha:]+\.\ ")
```

```
## [1] "Claibourn" "Ford"      "Draber"
```

- ▶ Replace "Group", "grp", and "group" with "G"

```
group <- c("Group 1", "grp 1", "group 3")  
str_replace(group, "(Group |grp |group )", "G")
```

```
## [1] "G1" "G1" "G3"
```

## Binding data frames

- ▶ Row binding: stacking data frames on top of one another
- ▶ Column binding: setting data frames next to each other
- ▶ Row binding is more common; often used when reading in multiple files of matching structure that need to be combined into one data frame
- ▶ Example: importing 10 Excel worksheets for years 2000 - 2009, and then combining into one data frame
- ▶ dplyr functions: `bind_rows` and `bind_cols`

## Row binding

```
bind_rows(data_frame_01, data_frame_02)
```

data\_frame\_01

id	height	weight
1	62	124
2	68	187
3	63	115

data\_frame\_02

id	height	weight
4	69	166
5	70	192
6	63	121

### Bind Rows

id	height	weight
1	62	124
2	68	187
3	63	115
4	69	166
5	70	192
6	63	121

## Column binding

`bind_cols(data_frame_01, data_frame_02)`<sup>1</sup>

data\_frame\_01

id	height	weight
1	62	124
2	68	187
3	63	115

data\_frame\_02

gender	rating
F	4
M	6
F	3

### Bind Columns

id	height	weight	gender	rating
1	62	124	F	4
2	68	187	M	6
3	63	115	F	3

---

<sup>1</sup>data frames must have the same number of rows.

# Merging or Joining data frames

- ▶ Two types of merges, or joins:
  - ▶ **Mutating join:** Join data frames based on a common column (or columns), called “keys”
  - ▶ **Filtering join:** keep rows in one data frame based on having (or not having) membership in another data frame
- ▶ Frequently used to combine two different sources of data.
- ▶ Example: merge subject demographic data with subject lab data
- ▶ dplyr functions: `inner_join`, `left_join`, `right_join`, `full_join`, `semi_join`, `anti_join`

## Inner Join

Retain only those rows with matching keys in both data sets.

```
inner_join(band_members,band_instruments,by="name")
```

band\_members

name	band
Mick	Stones
John	Beatles
Paul	Beatles

band\_instruments

name	plays
John	guitar
Paul	bass
Keith	guitar

## Inner join

name	band	plays
John	Beatles	guitar
Paul	Beatles	bass

## Left Join

Retain everything in the left data set, merge matches from right.

```
left_join(band_members, band_instruments, by="name")
```

band\_members

name	band
Mick	Stones
John	Beatles
Paul	Beatles

band\_instruments

name	plays
John	guitar
Paul	bass
Keith	guitar

## Left join

name	band	plays
Mick	Stones	NA
John	Beatles	guitar
Paul	Beatles	bass

## Right Join

Retain everything in the right data set, merge from left.

```
right_join(band_members, band_instruments, by="name")
```

band\_members

name	band
Mick	Stones
John	Beatles
Paul	Beatles

band\_instruments

name	plays
John	guitar
Paul	bass
Keith	guitar

## Right join

name	band	plays
John	Beatles	guitar
Paul	Beatles	bass
Keith	NA	guitar



## Full Join

Retain all rows in both data sets.

```
full_join(band_members, band_instruments, by="name")
```

band\_members

name	band
Mick	Stones
John	Beatles
Paul	Beatles

band\_instruments

name	plays
John	guitar
Paul	bass
Keith	guitar

### Full join

name	band	plays
Mick	Stones	NA
John	Beatles	guitar
Paul	Beatles	bass
Keith	NA	guitar

## Semi Join

Retain all rows from left that have matching values in right.

```
semi_join(band_members, band_instruments, by="name")
```

band\_members

name	band
Mick	Stones
John	Beatles
Paul	Beatles

band\_instruments

name	plays
John	guitar
Paul	bass
Keith	guitar

### Semi join

name	band
John	Beatles
Paul	Beatles

## Anti Join

Retain all rows from left that do NOT have matching values in right.

`anti_join(band_members, band_instruments, by="name")`

band\_members

name	band
Mick	Stones
John	Beatles
Paul	Beatles

band\_instruments

name	plays
John	guitar
Paul	bass
Keith	guitar

### Anti join

name	band
Mick	Stones

# Reshaping data frames

- ▶ Taking column names and making them values in a single column, and vice versa; similar to transposing data in Excel
- ▶ Often expressed as reshaping “wide to long”, or “long to wide”
- ▶ Reshaping wide to long is very common in R; often needed to accommodate modeling and plotting functions
- ▶ `tidyr` functions: `pivot_longer` (wide to long) and `pivot_wider` (long to wide)

## Reshaping wide to long

Reshape columns week1 - week3 into two columns, one for the column headers (names\_to) and the other for the values in the columns (values\_to).

```
pivot_longer(wide_df, week1:week3,  
             names_to = "week", values_to = "count")
```

### Reshape wide to long

wide\_df

id	week1	week2	week3
1	34	35	42
2	23	27	29

long\_df

id	week	count
1	week1	34
1	week2	35
1	week3	42
2	week1	23
2	week2	27
2	week3	29

## References

**R for Data Science:** <http://r4ds.had.co.nz/>

Free online edition of the O'Reilly book *R for Data Science*, by Garrett Grolemund and Hadley Wickham.

The UVa Library also has a physical copy.

## See also

This workshop was previously offered as two workshops, where each went into more detail. Here are links to the materials.

Part 1: bind, merge, reshape

[http://bit.ly/dwr\\_01](http://bit.ly/dwr_01)

Part 2: dates and strings

[http://bit.ly/dwr\\_02](http://bit.ly/dwr_02)

# Thanks for coming

- ▶ For statistical consulting: [statlab@virginia.edu](mailto:statlab@virginia.edu)
- ▶ Sign up for more workshops or see past workshops:  
<http://data.library.virginia.edu/training/>
- ▶ Register for the Research Data Services newsletter to be notified of new workshops:  
<http://data.library.virginia.edu/newsletters/>